

Twitter Sentiment Analysis

Install packages

```
#Load libraries
library('quanteda')

## Package version: 3.2.0
## Unicode version: 13.0
## ICU version: 67.1

## Parallel computing: 8 of 8 threads used.
## See https://quanteda.io for tutorials and examples.
library('readtext')
library('tidyverse')

## -- Attaching packages ----- tidyverse 1.3.1 --
## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.4      v dplyr  1.0.7
## v tidyr   1.1.4      v stringr 1.4.0
## v readr   2.0.2      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
library('quanteda.textstats')
library("rtweet")

##
## Attaching package: 'rtweet'

## The following object is masked from 'package:purrr':
##
##     flatten

# #library("tidytext")
library("textstem")

## Loading required package: koRpus.lang.en
## Loading required package: koRpus
## Loading required package: sylly
## For information on available language packages for 'koRpus', run
##
##     available.koRpus.lang()
##
## and see ?install.koRpus.lang()
##
```

```
## Attaching package: 'koRpus'

## The following object is masked from 'package:readr':
##
##      tokenize

## The following objects are masked from 'package:quanteda':
##
##      tokens, types

library("httr")
library("jsonlite")
```

```
##
## Attaching package: 'jsonlite'

## The following object is masked from 'package:rtweet':
##
##      flatten

## The following object is masked from 'package:purrr':
##
##      flatten

library("dplyr")
library("wordcloud")
```

```
## Loading required package: RColorBrewer
```

Bringing in Tweets

Tweets have been saved as a csv and are imported here

```
# Reading in the saved tweets
swu_18000 <- as.data.frame(readr::read_csv("/Users/garethmoen/Documents/Data Science/Portfolio/Sentimen

## Warning: One or more parsing issues, see `problems()` for details
## Rows: 16961 Columns: 90

## -- Column specification -----
## Delimiter: ","
## chr  (27): screen_name, text, source, reply_to_screen_name, lang, quoted_tex...
## dbl  (19): user_id, status_id, display_text_width, reply_to_status_id, reply...
## lgl  (41): is_quote, is_retweet, quote_count, reply_count, hashtags, symbols...
## dtm   (3): created_at, quoted_created_at, account_created_at

##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.

# Original method used to bring in tweets
# Getting twitter trends for Ireland
# ire_trends <- get_trends(woeid = "Ireland")
# head(ire_trends)
# ire_trends_7mar <- get_trends(woeid = "23424803")

# Getting twitter trends for the world
# wor_trends <- get_trends(woeid = "World", lang = "en")
# Bringing in 18000 tweets with the hashtag #StandWithUkraine, tweets in English (5th Mar)
```

```
# swu_18000 <- search_tweets("#StandWithUkraine", n = 18000, include_rts = FALSE, type = "mixed", lang = "en")
# readr::write_csv(swu_18000, "swu_18000.csv") # Saving the results to file
```

Full project used 400,000 tweets but is too big for Github storage

```
# Bringing in 400,000 tweets with the hashtag #StandWithUkraine, tweets in English (5th Mar)
# swu_400000 <- search_tweets("#StandWithUkraine", n = 400000, include_rts = FALSE, type = "mixed", retweets = FALSE)
# swu_400000 <- subset(swu_400000, is.na(swu_400000$reply_to_status_id))
# readr::write_csv(swu_400000, "swu_400000.csv") # Saving the results to file
# swu_400000 <- readr::read_csv("swu_400000.csv") # Now the file is save, only read it in
```

Cleaning the tweets

```
# Tweets should be original tweets, not retweets, as the original API request filtered them out
# Selecting columns
df <- swu_18000 %>%
  select(status_id,
         created_at,
         text,
         source,
         screen_name, # used to remove irrelevant tweets later
         favorite_count,
         retweet_count
  ) # %>%
  #as.data.frame()

# Remove large swu_18000 tibble
#rm(swu_18000)

# Removal of irrelevant sources
df <- df[-which(df$screen_name == "ArvadaRadio"),] # radio station
df <- df[-which(df$screen_name == "EstellaBell1"),] # other news
df <- df[-which(df$screen_name == "AlenaKazakevich"),] # repeated tweets
df <- df[-which(df$screen_name == "Fidget02"),] # only hashtags

# Remove duplicated status IDs if needed
# df <- df[-which(duplicated(df$status_id)),]

# Filter for popular tweets with 1 or more in the 'favourite_count'
df <- df[df$favorite_count >= 1,]

# Randomise the order of the tweets
set.seed(1234)
rows <- sample(nrow(df))
df_shuf <- df[rows, ]

# Looking at tweets related to Putin
#df_putin <- dplyr::filter(df_shuf, grepl('\\bputin\\b', text) | grepl('\\bPutin\\b', text))

# Looking at tweets related to 'ukraine' or 'ukrainians'
df_ukr <- dplyr::filter(df_shuf, grepl('\\bukraine\\b', text) | grepl('\\bukrainians\\b', text) | grepl('\\bUkraine\\b', text) | grepl('\\bUkrainians\\b', text))
```

```

# Conversion to corpus
corp <- corpus(df_ukr,
              docid_field = "status_id",
              text_field = "text",
              unique_docnames = TRUE)

# summary(corp)

prep_toks <- function(text_corpus){
  toks <- quanteda::tokens(text_corpus,
                          include_docvars = TRUE) %>%
    tokens_tolower() %>%
    tokens_remove(stopwords("english"), padding = TRUE) %>%
    tokens_remove('[\\p{P}\\p{S}]', valuetype = 'regex', padding = TRUE)
  return(toks)
}

toks <- corp %>%
  prep_toks()

# head(toks)

```

Creating collocations

```

get_coll <- function(tokens){
  unsup_col <- textstat_collocations(tokens,
                                    method = "lambda",
                                    size = 2,
                                    min_count = 5,
                                    smoothing = 0.5)

  unsup_col <- unsup_col[order(-unsup_col$count),] # sort detected collocations by count (descending)
  return(unsup_col)
}

collocations <- get_coll(toks) # create collocations

toks <- tokens_compound(toks, pattern = collocations[collocations$z > 5]) # merge collocations into tokens

toks <- tokens_remove(toks, c("amp", "come", "months", "weeks", "analysts_said", "can", "today", "now",

toks <- quanteda::tokens(toks,
                        remove_numbers = TRUE,
                        remove_punct = TRUE,
                        remove_symbols = TRUE,
                        remove_hyphens = TRUE,
                        remove_separators = TRUE,
                        remove_url = TRUE) # remove other uninformative text

```

Warning: remove_hyphens argument is not used.

Creating collocations and returning them back to the main corpus may not be useful as some of the words with sentiment may be lost. So for example, ‘aggressor’ would be considered a negative sentiment, but ‘russian aggressor’ may not be considered such. So it’s best to leave the words as they are for the purpose of sentiment analysis.

```

# Creating a dfm
dfm <- dfm(toks) # create DFM
dfm <- dfm_trim(dfm, min_docfreq = 20) # trim DFM
#dfm <- dfm_tfidf(dfm) # weight DFM

# Top 20 most frequent words / expressions
dfm_plot <- textstat_frequency(dfm, n = 20)
# use textstat_frequency(dfm, n = 30, force = TRUE) if the DFM is already weighted
#dfm_plot <- dfm_plot[-1,]
dfm_plot

```

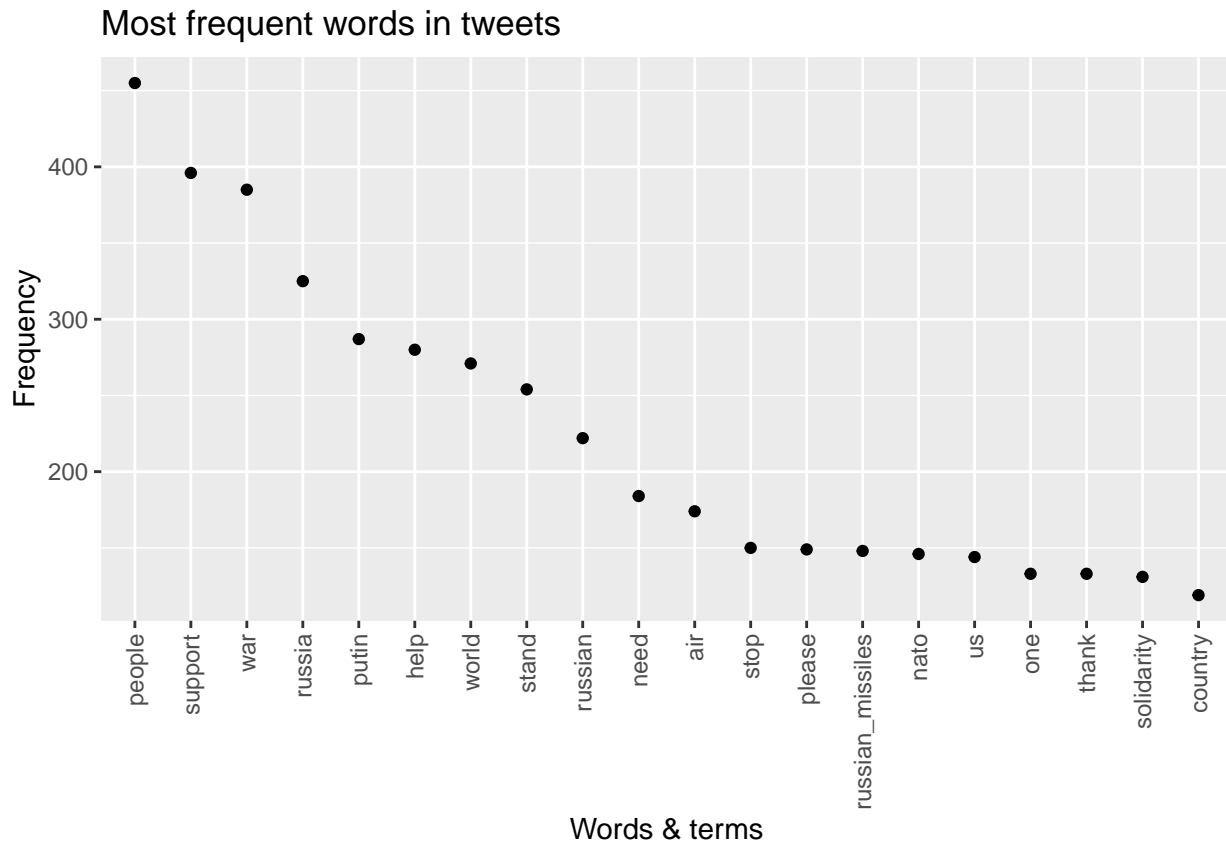
##	feature	frequency	rank	docfreq	group
## 1	people	455	1	411	all
## 2	support	396	2	370	all
## 3	war	385	3	373	all
## 4	russia	325	4	308	all
## 5	putin	287	5	274	all
## 6	help	280	6	265	all
## 7	world	271	7	247	all
## 8	stand	254	8	246	all
## 9	russian	222	9	215	all
## 10	need	184	10	179	all
## 11	air	174	11	137	all
## 12	stop	150	12	138	all
## 13	please	149	13	142	all
## 14	russian_missiles	148	14	127	all
## 15	nato	146	15	136	all
## 16	us	144	16	137	all
## 17	thank	133	17	117	all
## 18	one	133	17	125	all
## 19	solidarity	131	19	131	all
## 20	country	119	20	117	all

Plot the most frequent words & expressions

```

dfm_plot %>%
  ggplot(aes(x = reorder(feature, -frequency), y = frequency)) +
  geom_point() +
  theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust = 1)) +
  labs(title = "Most frequent words in tweets",
       x = "Words & terms",
       y = "Frequency")

```



Adding sentiment to the data

```
# General sentiment
dfm_sentiment <- dfm_lookup(dfm(toks), data_dictionary_LSD2015) %>%
  dfm_group(groups = dfm@docvars$created_at)
# dfm_sentiment <- dfm_lookup(dfm, data_dictionary_LSD2015) %>%
#   dfm_group(groups = date)
```

Plot of sentiment over time

```
# Only keep column 1 - 2
dfm_sentiment <- dfm_lookup(dfm(toks), data_dictionary_LSD2015[1:2])
# Next prepare to plot
docvars(dfm_sentiment, "prop_negative") <- as.numeric(dfm_sentiment[,1] / ntoken(dfm_sentiment))
docvars(dfm_sentiment, "prop_positive") <- as.numeric(dfm_sentiment[,2] / ntoken(dfm_sentiment))


docvars(dfm_sentiment, "net_sentiment") <- docvars(dfm_sentiment, "prop_positive") - docvars(dfm_sentiment, "prop_negative")

docvars(dfm_sentiment) %>%
  ggplot(aes(x = dfm@docvars$created_at, y = net_sentiment)) +
  geom_smooth() +
  labs(title = "Tweets with 'Ukraine', sentiment analysis over time",
       x = "Date",
       y = "Sentiment")
```

```
## `geom_smooth()` using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

```
## Warning: Removed 668 rows containing non-finite values (stat_smooth).
```

Tweets with 'Ukraine', sentiment analysis over time



Wordcloud of most frequent words

```
dfwc <- as.data.frame(dfm_plot)
wordcloud(words = dfwc$feature, freq = dfwc$frequency, min.freq = 1, max.words=150, random.order=FALSE,
```

